



**Министерство науки и высшего образования
Российской Федерации Федеральное
государственное бюджетное образовательное
учреждение высшего образования «Московский
государственный технический университет имени Н.Э.
Баумана
(национальный исследовательский университет)»
(МГТУ им. Н.Э. Баумана)**

**Факультет «Информатика и системы управления»
Кафедра ИУ5 «Системы обработки информации и управления»**

Курс «Технологии машинного обучения»

Лабораторная работа № 2

Выполнил:
студент группы ИУ5-63Б
Попов С. Д.

Проверил:
Гапанюк Ю. Е.

2024 г.

Задание

Выбрать набор данных (датасет), содержащий категориальные признаки и пропуски в данных. Для выполнения следующих пунктов можно использовать несколько различных наборов данных (один для обработки пропусков, другой для категориальных признаков и т.д.) Для выбранного датасета (датасетов) на основе материалов лекции решить следующие задачи:

- обработку пропусков в данных;
- кодирование категориальных признаков;
- масштабирование данных.

Датасет Melbourne Housing Snapshot

<https://www.kaggle.com/datasets/dansbecker/melbourne-housing-snapshot?resource=download>

```
In [2]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style="ticks")
```

```
In [3]: from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.impute import MissingIndicator
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OrdinalEncoder
```

```
In [4]: data = pd.read_csv('/kaggle/input/melb-data/melb_data.csv', sep=",")
```

```
In [5]: data.shape
```

```
Out[5]: (13580, 21)
```

```
In [6]: data.dtypes
```

```
Out [6]: Suburb          object
Address         object
Rooms           int64
Type            object
Price           float64
Method          object
SellerG         object
Date            object
Distance        float64
Postcode        float64
Bedroom2        float64
Bathroom        float64
Car             float64
Landsize        float64
BuildingArea    float64
YearBuilt       float64
CouncilArea     object
Lattitude       float64
Longitude       float64
Regionname      object
Propertycount   float64
dtype: object
```

```
In [7]: data.isnull().sum()
```

```
Out [7]: Suburb          0
Address         0
Rooms           0
Type            0
Price           0
Method          0
SellerG         0
Date            0
Distance        0
Postcode        0
Bedroom2        0
Bathroom        0
Car             62
Landsize        0
BuildingArea    6450
YearBuilt       5375
CouncilArea     1369
Lattitude       0
Longitude       0
Regionname      0
Propertycount   0
dtype: int64
```

```
In [8]: data.head()
```

Out [8]:		Suburb	Address	Rooms	Type	Price	Method	SellerG	Date	Distance	Postc
0	Abbotsford	85 Turner St		2	h	1480000.0	S	Biggin	3/12/2016	2.5	30
1	Abbotsford	25 Bloomburg St		2	h	1035000.0	S	Biggin	4/02/2016	2.5	30
2	Abbotsford	5 Charles St		3	h	1465000.0	SP	Biggin	4/03/2017	2.5	30
3	Abbotsford	40 Federation La		3	h	850000.0	PI	Biggin	4/03/2017	2.5	30
4	Abbotsford	55a Park St		4	h	1600000.0	VB	Nelson	4/06/2016	2.5	30

5 rows x 21 columns

Обработка пропусков категориальных признаков

```
In [9]: total_count = data.shape[0]

cat_cols = []
for col in data.columns:
    # Количество пустых значений
    temp_null_count = data[data[col].isnull()].shape[0]
    dt = str(data[col].dtype)
    if temp_null_count > 0 and (dt == 'object'):
        cat_cols.append(col)
        temp_perc = round((temp_null_count / total_count) * 100.0, 2)
        print('Колонка {}. Тип данных {}. Количество пустых значений {}, {}%.'.format
```

Колонка CouncilArea. Тип данных object. Количество пустых значений 1369, 10.08%.

```
In [10]: data['CouncilArea'].describe()
```

```
Out[10]: count      12211
unique         33
top      Moreland
freq         1163
Name: CouncilArea, dtype: object
```

```
In [11]: data['CouncilArea'].unique()
```

```
Out[11]: array(['Yarra', 'Moonee Valley', 'Port Phillip', 'Darebin', 'Hobsons Bay',
                'Stonnington', 'Boroondara', 'Monash', 'Glen Eira', 'Whitehorse',
                'Maribyrnong', 'Bayside', 'Moreland', 'Manningham', 'Banyule',
                'Melbourne', 'Kingston', 'Brimbank', 'Hume', nan, 'Knox',
                'Maroondah', 'Casey', 'Melton', 'Greater Dandenong', 'Nillumbik',
                'Whittlesea', 'Frankston', 'Macedon Ranges', 'Yarra Ranges',
                'Wyndham', 'Cardinia', 'Unavailable', 'Moorabool'], dtype=object)
```

```
In [12]: # Проверим корреляцию региона и цены
# Для этого закодируем колонку 'CouncilArea' в зависимости от ее средней цены
df_ca_and_target = data[['Price', 'CouncilArea']].copy()
df_ca_and_target.dropna(subset=['CouncilArea'], inplace=True)

encoder_map = {}
for index, row in df_ca_and_target.groupby(by='CouncilArea').mean().sort_values(by='Price').iteritems():
    encoder_map[row['CouncilArea']] = index
    # print(index, row)
```

```
le = LabelEncoder()
df_ca_and_target['CouncilArea'] = df_ca_and_target['CouncilArea'].map(encoder_map)
df_ca_and_target.corr()
```

```
Out[12]:
```

	Price	CouncilArea
Price	1.000000	0.423142
CouncilArea	0.423142	1.000000

```
In [13]: # У нас есть небольшая корреляция между целевым признаком и ценой, поэтому просто уда
data.dropna(subset=['CouncilArea'], inplace=True)
```

```
In [14]: data['CouncilArea'].isnull().sum()
```

```
Out[14]: 0
```

Обработка пропусков числовых признаков

```
In [15]: num_cols = []
for col in data.columns:
    # Количество пустых значений
    temp_null_count = data[data[col].isnull()].shape[0]
    dt = str(data[col].dtype)
    if temp_null_count > 0 and (dt == 'float64' or dt == 'int64'):
        num_cols.append(col)
        temp_perc = round((temp_null_count / total_count) * 100.0, 2)
        print('Колонка {}. Тип данных {}. Количество пустых значений {}, {}%.'.format
```

Колонка BuildingArea. Тип данных float64. Количество пустых значений 5765, 42.45%.
Колонка YearBuilt. Тип данных float64. Количество пустых значений 4763, 35.07%.

```
In [16]: # Заполним пропуски в числовых признаках медианой, тк она меньше всего подвержена выб

def fill_missing_median(dataset, column):
    temp_data = dataset[[column]]

    imp_num = SimpleImputer(strategy='median')
    data_num_imp = imp_num.fit_transform(temp_data)

    return data_num_imp
```

```
In [17]: cols_for_imp = ['BuildingArea', 'YearBuilt']
data[cols_for_imp].isnull().sum()
```

```
Out[17]: BuildingArea    5765
YearBuilt      4763
dtype: int64
```

```
In [18]: for col in cols_for_imp:
    data[col] = fill_missing_median(data, col)
    print(f'After {col} imp:')
    print(data[cols_for_imp].isnull().sum())
    print()
```

```
After BuildingArea imp:
BuildingArea      0
YearBuilt         4763
dtype: int64
```

```
After YearBuilt imp:
BuildingArea      0
YearBuilt         0
dtype: int64
```

```
In [19]: data.isnull().sum()
```

```
Out[19]: Suburb      0
Address    0
Rooms      0
Type       0
Price      0
Method     0
SellerG    0
Date       0
Distance   0
Postcode   0
Bedroom2   0
Bathroom   0
Car        0
Landsize   0
BuildingArea 0
YearBuilt   0
CouncilArea 0
Latitude   0
Longitude  0
Regionname  0
Propertycount 0
dtype: int64
```

Преобразование категориальных признаков в числовые

```
In [20]: data.dtypes.loc[lambda x: x == 'object']
```

```
Out[20]: Suburb      object
Address    object
Type       object
Method     object
SellerG    object
Date       object
CouncilArea object
Regionname object
dtype: object
```

Date

```
In [21]: data['Date'].unique()
```

```
Out [21]: array(['3/12/2016', '4/02/2016', '4/03/2017', '4/06/2016', '7/05/2016',
               '8/10/2016', '12/11/2016', '13/08/2016', '14/05/2016',
               '15/10/2016', '16/07/2016', '17/09/2016', '18/03/2017',
               '18/06/2016', '19/11/2016', '22/08/2016', '24/09/2016',
               '25/02/2017', '26/07/2016', '27/11/2016', '28/05/2016',
               '30/07/2016', '3/09/2016', '6/08/2016', '7/11/2016', '10/09/2016',
               '10/12/2016', '11/02/2017', '16/04/2016', '22/05/2016',
               '23/04/2016', '12/06/2016', '27/06/2016', '28/08/2016',
               '11/03/2017', '28/01/2016', '6/05/2017', '8/04/2017', '29/04/2017',
               '13/05/2017', '20/05/2017', '22/04/2017', '1/07/2017', '3/06/2017',
               '17/06/2017', '24/06/2017', '27/05/2017', '8/07/2017',
               '12/08/2017', '15/07/2017', '22/07/2017', '29/07/2017'],
              dtype=object)
```

```
In [22]: # Разделим поле Date на day, month, year
date = pd.to_datetime(data['Date'], format='mixed')

data['Day'] = date.dt.day
data['Month'] = date.dt.month
data['Year'] = date.dt.year
data.drop(columns='Date', inplace=True)

data[['Day', 'Month', 'Year']]
```

```
Out [22]:
```

	Day	Month	Year
0	12	3	2016
1	2	4	2016
2	3	4	2017
3	3	4	2017
4	6	4	2016
...
12208	29	7	2017
12209	29	7	2017
12210	29	7	2017
12211	29	7	2017
12212	29	7	2017

12211 rows x 3 columns

Adress

```
In [23]: # Преобразуем улицу в адрес

data['Street'] = data['Address'].map(lambda addr: addr.split()[1])
data.drop(columns='Address', inplace=True)

data['Street']
```

```
Out [23]: 0      Turner
          1      Bloomburg
          2      Charles
          3      Federation
          4      Park
          ...
          12208    Pasco
          12209    Peel
          12210    Saltlake
          12211    Adeney
          12212    Pentland
          Name: Street, Length: 12211, dtype: object
```

```
In [24]: data['Street'].unique().shape
```

```
Out [24]: (3721,)
```

Other

```
In [25]: data.dtypes.loc[lambda x: x == 'object']
```

```
Out [25]: Suburb      object
          Type        object
          Method      object
          SellerG     object
          CouncilArea  object
          Regionname   object
          Street      object
          dtype: object
```

```
In [26]: obj_cols = list(data.dtypes.loc[lambda x: x == 'object'].index)
          obj_cols
```

```
Out [26]: ['Suburb', 'Type', 'Method', 'SellerG', 'CouncilArea', 'Regionname', 'Street']
```

```
In [27]: data_obj = data[obj_cols]

          oe = OrdinalEncoder()
          data_obj_enc = oe.fit_transform(data_obj)
```

```
In [28]: data[obj_cols] = data_obj_enc
          data[obj_cols]
```


Out [28]:

	Suburb	Type	Method	SellerG	CouncilArea	Regionname	Street
0	0.0	0.0	1.0	22.0	31.0	2.0	3377.0
1	0.0	0.0	1.0	22.0	31.0	2.0	377.0
2	0.0	0.0	3.0	22.0	31.0	2.0	668.0
3	0.0	0.0	0.0	22.0	31.0	2.0	1246.0
4	0.0	0.0	4.0	146.0	31.0	2.0	2657.0
...
12208	299.0	0.0	1.0	100.0	10.0	6.0	2670.0
12209	301.0	2.0	0.0	245.0	26.0	5.0	2691.0
12210	302.0	0.0	1.0	207.0	29.0	2.0	2995.0
12211	307.0	0.0	3.0	245.0	16.0	6.0	24.0
12212	307.0	0.0	4.0	224.0	16.0	6.0	2703.0

12211 rows × 7 columns

In [29]: data.describe()

Out [29]:

	Suburb	Rooms	Type	Price	Method	SellerG	
count	12211.000000	12211.000000	12211.000000	1.221100e+04	12211.000000	12211.000000	122
mean	152.520514	2.894194	0.570715	1.063692e+06	1.404389	123.668741	
std	89.000246	0.959341	0.854515	6.388613e+05	1.117471	73.079999	
min	0.000000	1.000000	0.000000	8.500000e+04	0.000000	0.000000	
25%	67.000000	2.000000	0.000000	6.400000e+05	1.000000	70.000000	
50%	153.000000	3.000000	0.000000	8.950000e+05	1.000000	128.000000	
75%	229.000000	3.000000	1.000000	1.320000e+06	1.000000	183.000000	
max	307.000000	10.000000	2.000000	9.000000e+06	4.000000	252.000000	

8 rows × 23 columns

In [30]: data.dtypes.loc[lambda x: x == 'object'].size


Out [30]: 0

Масштабирование данных


До применения Z скеллинга



```
for col in data.columns:
    plt.hist(data[col], 100)
    plt.title(col)
    plt.show()
```

Suburb




Rooms







Method




SellerG




Distance




Postcode





Bedroom2




Bathroom








BuildingArea




YearBuilt










Longitude

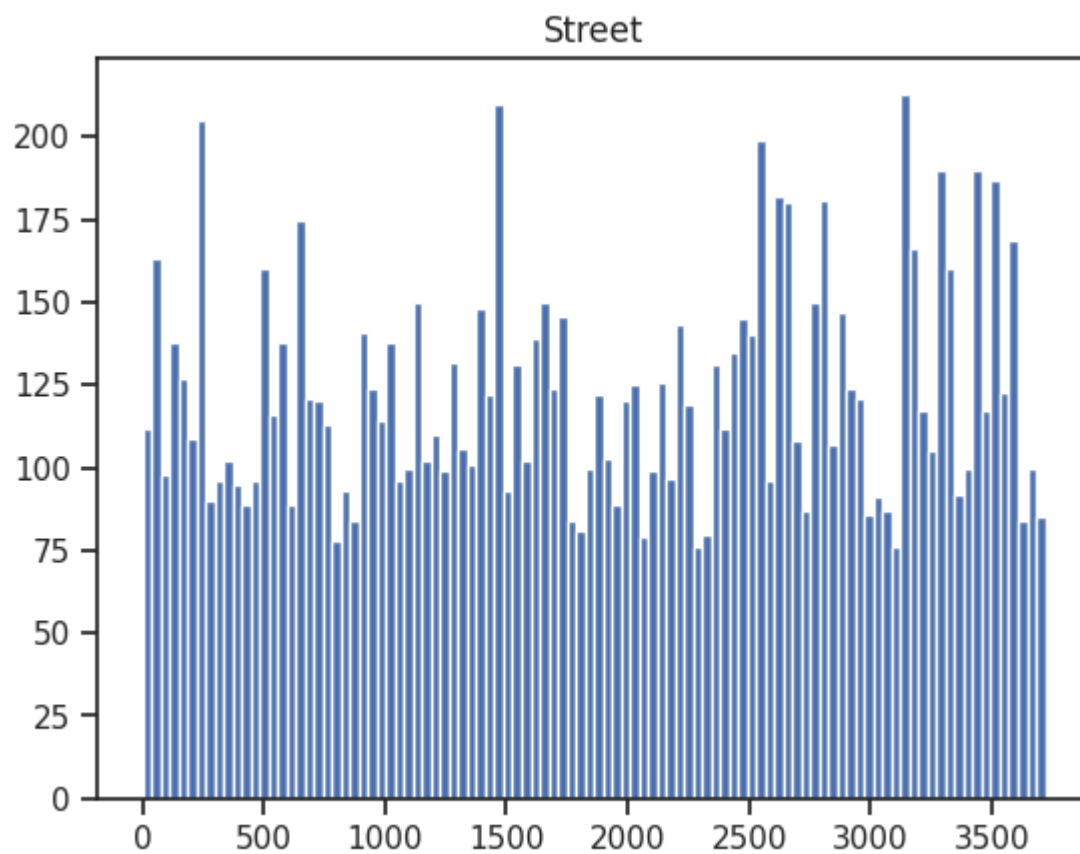


Regionname









После


```
In [32]: from sklearn.preprocessing import MinMaxScaler, StandardScaler, Normalizer
```

```
In [33]: sc2 = StandardScaler()  
sc2_data = sc2.fit_transform(data)
```


```
In [34]: sc2_data = pd.DataFrame(sc2_data, columns=data.columns)
```

```
In [35]: # После  
for col in data.columns:  
    plt.hist(sc2_data[col], 100)  
    plt.title(col)  
    plt.show()
```


Suburb




Rooms




Type




Price




Method




SellerG




Distance




Postcode





Bedroom2




Bathroom








BuildingArea




YearBuilt










Longitude




Regionname









In []:

Обучение с фиксированным гиперпараметром

```
In [36]: from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsRegressor, KNeighborsClassifier
from sklearn.model_selection import learning_curve, validation_curve
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score, explain
```

```
In [37]: data = sc2_data
target = sc2_data['Price']

data.drop(columns=['Price'], inplace=True)
```

```
In [38]: data.shape
```

```
Out[38]: (12211, 22)
```

```
In [39]: data.describe()
```

Out [39]:

	Suburb	Rooms	Type	Method	SellerG	Distance
count	1.221100e+04	1.221100e+04	1.221100e+04	1.221100e+04	1.221100e+04	1.221100e+04
mean	-1.675836e-16	9.775709e-17	-5.586119e-17	-8.728311e-17	6.749894e-17	4.096487e-16
std	1.000041e+00	1.000041e+00	1.000041e+00	1.000041e+00	1.000041e+00	1.000041e+00
min	-1.713779e+00	-1.974554e+00	-6.679092e-01	-1.256809e+00	-1.692307e+00	-1.772836e+00
25%	-9.609414e-01	-9.321296e-01	-6.679092e-01	-3.618941e-01	-7.344135e-01	-7.016215e-01
50%	5.387686e-03	1.102950e-01	-6.679092e-01	-3.618941e-01	5.926979e-02	-1.024677e-01
75%	8.593529e-01	1.102950e-01	5.023935e-01	-3.618941e-01	8.119005e-01	5.148423e-01
max	1.735791e+00	7.407267e+00	1.672696e+00	2.322850e+00	1.756110e+00	6.833191e+00

8 rows x 22 columns

In [40]:

target.describe()

Out [40]:

count 1.221100e+04
mean -1.629285e-16
std 1.000041e+00
min -1.531994e+00
25% -6.632256e-01
50% -2.640617e-01
75% 4.012117e-01
max 1.242309e+01
Name: Price, dtype: float64

In [41]:

X_train, X_test, y_train, y_test = train_test_split(
 data, target, test_size=0.5, random_state=1)

In [42]:

K = 70

cl = KNeighborsRegressor(n_neighbors=K)
cl.fit(X_train, y_train)
target_train = cl.predict(X_train)
target_test = cl.predict(X_test)

In [43]:

def scorer(y_true, y_pred):
 scores = {
 'r2': r2_score,
 'mean_squared_error': mean_squared_error,
 'mean_absolute_error': mean_absolute_error,
 'r2_score': r2_score,
 'explained_variance_score': explained_variance_score,
 'mean_pinball_loss': mean_pinball_loss,
 'd2_pinball_score': d2_pinball_score,
 'd2_absolute_error_score': d2_absolute_error_score
 }

 for score, score_func in scores.items():
 scores[score] = score_func(y_true, y_pred)
 return scores

In [44]:

print('Train: ', scorer(y_train, target_train))
print()
print('Test :', scorer(y_test, target_test))

```
Train: {'r2': 0.6074478363288117, 'mean_squared_error': 0.4054128638342073, 'mean_absolute_error': 0.38465733025879834, 'r2_score': 0.6074478363288117, 'explained_variance_score': 0.6080485130705291, 'mean_pinball_loss': 0.19232866512939917, 'd2_pinball_score': 0.4454301531494722, 'd2_absolute_error_score': 0.4454301531494722}
```

```
Test : {'r2': 0.6358064659388151, 'mean_squared_error': 0.3521518459937633, 'mean_absolute_error': 0.3775204179759272, 'r2_score': 0.6358064659388151, 'explained_variance_score': 0.6359862432962011, 'mean_pinball_loss': 0.1887602089879636, 'd2_pinball_score': 0.4426146306868046, 'd2_absolute_error_score': 0.4426146306868046}
```

Поиск лучших гиперпараметров

```
In [45]: from sklearn.model_selection import KFold, RepeatedKFold, LeaveOneOut, LeavePOut, ShuffleSplit
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
```

Вспомогательные функции

```
In [46]: def plot_validation_curve(estimator, title, X, y,
                                   param_name, param_range, cv,
                                   scoring='accuracy'):

    train_scores, test_scores = validation_curve(
        estimator, X, y, param_name=param_name, param_range=param_range,
        cv=cv, scoring=scoring, n_jobs=1)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)

    plt.title(title)
    plt.xlabel(param_name)
    plt.ylabel(str(scoring))
    # plt.ylim(0.0, 1.1)
    lw = 2
    plt.plot(param_range, train_scores_mean, label="Training score",
             color="darkorange", lw=lw)
    plt.fill_between(param_range, train_scores_mean - train_scores_std,
                    train_scores_mean + train_scores_std, alpha=0.4,
                    color="darkorange", lw=lw)
    plt.plot(param_range, test_scores_mean, label="Cross-validation score",
             color="navy", lw=lw)
    plt.fill_between(param_range, test_scores_mean - test_scores_std,
                    test_scores_mean + test_scores_std, alpha=0.2,
                    color="navy", lw=lw)
    plt.legend(loc="best")
    return plt
```

```
In [47]: def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                                   n_jobs=None, train_sizes=np.linspace(.1, 1.0, 5), scoring='accuracy'):
    """
    Generate a simple plot of the test and training learning curve.

    Parameters
    -----
    estimator : object type that implements the "fit" and "predict" methods
        An object of that type which is cloned for each validation.

    title : string
        Title for the chart.

    X : array-like, shape (n_samples, n_features)
        Training vector, where n_samples is the number of samples and
        n_features is the number of features.
```

`y` : array-like, shape (n_samples) or (n_samples, n_features), optional
Target relative to X for classification or regression;
None for unsupervised learning.

`ylim` : tuple, shape (ymin, ymax), optional
Defines minimum and maximum yvalues plotted.

`cv` : int, cross-validation generator or an iterable, optional
Determines the cross-validation splitting strategy.
Possible inputs for cv are:

- None, to use the default 3-fold cross-validation,
- integer, to specify the number of folds.
- :term:`CV splitter`,
- An iterable yielding (train, test) splits as arrays of indices.

For integer/None inputs, if ``y`` is binary or multiclass,
:class:`StratifiedKFold` used. If the estimator is not a classifier
or if ``y`` is neither binary nor multiclass, :class:`KFold` is used.

Refer :ref:`User Guide <cross_validation>` for the various
cross-validators that can be used here.

`n_jobs` : int or None, optional (default=None)
Number of jobs to run in parallel.
``None`` means 1 unless in a :obj:`joblib.parallel_backend` context.
``-1`` means using all processors. See :term:`Glossary <n_jobs>`
for more details.

`train_sizes` : array-like, shape (n_ticks,), dtype float or int
Relative or absolute numbers of training examples that will be used to
generate the learning curve. If the dtype is float, it is regarded as a
fraction of the maximum size of the training set (that is determined
by the selected validation method), i.e. it has to be within (0, 1].
Otherwise it is interpreted as absolute sizes of the training sets.
Note that for classification the number of samples usually have to
be big enough to contain at least one sample from each class.
(default: np.linspace(0.1, 1.0, 5))

```
"""
plt.figure()
plt.title(title)
if ylim is not None:
    plt.ylim(*ylim)
plt.xlabel("Training examples")
plt.ylabel(scoring)
train_sizes, train_scores, test_scores = learning_curve(
    estimator, X, y, cv=cv, scoring=scoring, n_jobs=n_jobs, train_sizes=train_siz
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)
plt.grid()

plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.3,
                 color="r")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.1, color="g")
plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")

plt.legend(loc="best")
return plt
```

```
In [48]: def get_best_param_for_score(results, score):
        return results['params'][np.where(results[f"rank_test_{score}"] == 1)[0][0]]

def get_best_param_for_scores(results, scores):
    return {score: get_best_param_for_score(results, score) for score in scores}

In [68]: def show_plots_for_scores(n_range, results, scores):
        for score in scores:
            plt.plot(n_range, results[f"mean_test_{score}"])
            plt.title(score)
            plt.show()
```

Поиск

```
In [50]: scores = [
        'r2',
        'neg_mean_squared_error',
        'neg_mean_absolute_error',
    ]
```

RandomizedSearch + KFold

```
In [51]: n_range_1 = np.array(range(3,100,10))
        tuned_parameters_1 = [{'n_neighbors': n_range_1}]
        tuned_parameters_1
```

```
Out[51]: [{'n_neighbors': array([ 3, 13, 23, 33, 43, 53, 63, 73, 83, 93])}]
```

```
In [52]: data.shape
```

```
Out[52]: (12211, 22)
```

```
In [53]: kf = KFold(n_splits=24) # 500 samples for 1 fold
```

```
In [54]: %%time
        clf_rs = RandomizedSearchCV(KNeighborsRegressor(), tuned_parameters_1, cv=kf, scoring='r2')
        s = clf_rs.fit(X_train, y_train)
```



CPU times: user 14.7 s, sys: 39.4 ms, total: 14.8 s


Wall time: 4.12 s

```
In [55]: get_best_param_for_scores(clf_rs.cv_results_, scores)
```

```
Out[55]: {'r2': {'n_neighbors': 13},
        'neg_mean_squared_error': {'n_neighbors': 13},
        'neg_mean_absolute_error': {'n_neighbors': 13}}
```

```
In [69]: show_plots_for_scores(n_range_1, clf_rs.cv_results_, scores)
```





GridSearch + LeavePOut folds

```
In [57]: n_range_2 = np.array(range(10,20,1))
tuned_parameters_2 = [{'n_neighbors': n_range_2}]
tuned_parameters_2
```

```
Out[57]: [{'n_neighbors': array([10, 11, 12, 13, 14, 15, 16, 17, 18, 19])}]
```

```
In [58]: data.shape[0] / 3
```

```
Out[58]: 4070.3333333333335
```

```
In [62]: rkf = RepeatedKFold(n_splits=12, n_repeats=2)
len(list(rkf.split(X_train, y_train)))
```

```
Out[62]: 24
```

```
In [63]: %time
clf_gs = GridSearchCV(KNeighborsRegressor(), tuned_parameters_2, cv=rkf, scoring=scor
clf_gs.fit(X_train, y_train)
```



CPU times: user 15.6 s, sys: 23.4 ms, total: 15.6 s
Wall time: 4.46 s


```
Out[63]: 
> GridSearchCV
> estimator: KNeighborsRegressor
  > KNeighborsRegressor
```

```
In [71]: get_best_param_for_scores(clf_gs.cv_results_, scores)
```

```
Out[71]: {'r2': {'n_neighbors': 11},
'neg_mean_squared_error': {'n_neighbors': 11},
'neg_mean_absolute_error': {'n_neighbors': 14}}
```

```
In [70]: show_plots_for_scores(n_range_2, clf_gs.cv_results_, scores)
```






Validation curve

```
In [76]: plot_validation_curve(KNeighborsRegressor(), 'knn',
                               X_test, y_test,
                               param_name='n_neighbors', param_range=n_range_1,
                               cv=12, scoring='neg_mean_squared_error')
```


```
Out[76]: <module 'matplotlib.pyplot' from '/opt/conda/lib/python3.10/site-packages/matplotliblib/pyplot.py'>
```



Learning Curve for best hyperparams

```
In [78]: K = 11
plot_learning_curve(KNeighborsRegressor(n_neighbors=11), 'n_neighbors=11',
                    X_train, y_train, cv=20, scoring='neg_mean_squared_error')
```

```
Out[78]: <module 'matplotlib.pyplot' from '/opt/conda/lib/python3.10/site-packages/matplotlib/
pyplot.py'>
```



```
In [79]: K = 14
plot_learning_curve(KNeighborsRegressor(n_neighbors=14), 'n_neighbors=11',
                    X_train, y_train, cv=20, scoring='neg_mean_absolute_error')
```

```
Out[79]: <module 'matplotlib.pyplot' from '/opt/conda/lib/python3.10/site-packages/matplotlib/
pyplot.py'>
```

n_neighbors=11

